

Multiple Linear Regression

It is a statistical method used to analyze the relationship between multiple independent variables and a single dependent variable. It is an extension of simple linear regression which only consists of one independent variable.

The relationship between the dependent variable and independent variables is modeled as a linear equation. The goal is to find the coefficients that best fit the data and allow us to predict the dependent variable based on the values of the independent variable.

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import math
5 import time
```

In [2]:

```
1 housing_data = pd.read_csv('Bangalore.csv')
2 housing_data.head()
```

Out[2]:

	Price	Area	Location	NumBedrooms	Resale	MaintenanceStaff	Gymnasium	SwimmingPool	LandscapedGardens	JoggingTrack	...	LiftAvailable	BED	V
0	30000000	3340	JP Nagar Phase 1	4	0	1	1	1	1	1	...	1	0	
1	7888000	1045	Dasarahalli on Tumkur Road	2	0	0	1	1	1	1	...	1	0	
2	4866000	1179	Kannur on Thanisandra Main Road	2	0	0	1	1	1	1	...	1	0	
3	8358000	1675	Doddanekundi	3	0	0	0	0	0	0	...	1	0	
4	6845000	1670	Kengeri	3	0	1	1	1	1	1	...	1	0	

5 rows × 40 columns

In [3]:

```
1 varlist = ['Price', 'Area', 'NumBedrooms', 'Resale', 'MaintenanceStaff', 'Gymnasium', 'SwimmingPool', 'LandscapedGardens', 'IndoorGames']
2 data = housing_data[varlist].copy()
```

In [4]:

```
1 # checking for null values
2 print(data.info())
3
4 # checking for outliers
5 print(data.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6207 entries, 0 to 6206
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Price                 6207 non-null  int64
1   Area                 6207 non-null  int64
2   NumBedrooms          6207 non-null  int64
3   Resale               6207 non-null  int64
4   MaintenanceStaff     6207 non-null  int64
5   Gymnasium            6207 non-null  int64
6   SwimmingPool         6207 non-null  int64
7   LandscapedGardens    6207 non-null  int64
8   IndoorGames          6207 non-null  int64
dtypes: int64(9)
memory usage: 436.5 KB
None
```

	Price	Area	NumBedrooms	Resale	MaintenanceStaff	\
count	6.207000e+03	6207.000000	6207.000000	6207.000000	6207.000000	
mean	1.058510e+07	1526.094248	2.556952	0.078782	6.208797	
std	1.410943e+07	764.845609	0.694300	0.269420	4.126883	
min	2.000000e+06	415.000000	1.000000	0.000000	0.000000	
25%	5.000000e+06	1110.000000	2.000000	0.000000	0.000000	
50%	7.368000e+06	1340.000000	3.000000	0.000000	9.000000	
75%	1.070000e+07	1662.500000	3.000000	0.000000	9.000000	
max	3.000000e+08	9900.000000	7.000000	1.000000	9.000000	


	Gymnasium	SwimmingPool	LandscapedGardens	IndoorGames
count	6207.000000	6207.000000	6207.000000	6207.000000
mean	6.461576	6.436121	6.382471	6.348639
std	3.752421	3.792567	3.875271	3.926154
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	1.000000	1.000000	1.000000
50%	9.000000	9.000000	9.000000	9.000000
75%	9.000000	9.000000	9.000000	9.000000
max	9.000000	9.000000	9.000000	9.000000

```

In [5]: 1 # some of the features like Gymnasium should have bool values and the mean should not be more than 1
2 # plotting the to check if scaling is required
3 fig,ax = plt.subplots(1,8, figsize=(25,4), sharey = True)
4 for i in range(len(ax)):
5     ax[i].scatter(data[varlist[i+1]],data[varlist[0]])
6     ax[i].set_xlabel(varlist[i+1])
7     ax[0].set_ylabel('Price')

```

Out[5]: Text(0, 0.5, 'Price')



```

In [6]: 1 # MaintenanceStaff, Gymnasium, SwimmingPool, LandscapedGardens and IndoorGames needs to be manipulated
2 MaintenanceStaff = [1 if i>1 else 0 for i in data.MaintenanceStaff]
3 Gymnasium = [1 if i>1 else 0 for i in data.Gymnasium]
4 SwimmingPool = [1 if i>1 else 0 for i in data.SwimmingPool]
5 LandscapedGardens = [1 if i>1 else 0 for i in data.LandscapedGardens]
6 IndoorGames = [1 if i>1 else 0 for i in data.IndoorGames]

```

```

In [7]: 1 # all the features are scaled except Area and NumBedrooms
2 # performing z-score normalization on Area
3 std_dev = np.std(data.Area)
4 mean = np.mean(data.Area)
5 Area = [(i-mean)/std_dev for i in data.Area]
6
7 # performing z-score normalization on NumBedrooms
8 std_dev = np.std(data.NumBedrooms)
9 mean = np.mean(data.NumBedrooms)
10 NumBedrooms = [(i-mean)/std_dev for i in data.NumBedrooms]


```

```

In [8]: 1 # plotting Price to check if scaling is required
2 plt.scatter(data.Price,list(range(len(data.Price))))


```

Out[8]: <matplotlib.collections.PathCollection at 0x120b7410>




```
In [9]: 1 # some values are too large
2 # performing z-score normalization on price
3 std_dev = np.std(data.Price)
4 mean = np.mean(data.Price)
5 Price = [(i-std_dev)/mean for i in data.Price]
6 plt.scatter(Price,list(range(len(Price))))
```

Out[9]: <matplotlib.collections.PathCollection at 0x1210ca30>



```
In [10]: 1 # combining scaled data into a single dataframe
2 scaled_data = pd.DataFrame()
3 for var in varlist:
4     if var == 'Resale':
5         scaled_data[var] = data.Resale
6     else:
7         scaled_data[var] = eval(var)
8
9 fig,ax = plt.subplots(1,8, figsize=(25,4), sharey = True)
10 for i in range(len(ax)):
11     ax[i].scatter(scaled_data[varlist[i+1]],scaled_data[varlist[0]])
12     ax[i].set_xlabel(varlist[i+1])
13 ax[0].set_ylabel('Price')
```

Out[10]: Text(0, 0.5, 'Price')



```
In [11]: 1 # splitting the data into training and testing data
2 scaled_data = scaled_data.sample(frac=1, random_state=42).reset_index(drop=True)
3 train_to_test_ratio = 0.7
4 size = int(train_to_test_ratio * len(scaled_data))
5
6 data_train = scaled_data[size:]
7 data_test = scaled_data[:size]
8
9 x_train = data_train[varlist[1:]].values
10 y_train = data_train[varlist[0]].values
11
12 x_test = data_test[varlist[1:]].values
13 y_test = data_test[varlist[0]].values
14
```

Prediction function

We need to derive a function $f(\vec{w}, b)$ where $f(\vec{w}, b) = \vec{w} * \vec{x} + b$

here, \vec{w} is an array of weights = $[w_1, w_2, w_3, \dots, w_n]$

and b is the bias

This function predicts the value of y (price) for a given \vec{x} where \vec{x} is the array of features = $[x_1, x_2, x_3, \dots, x_n]$

```
In [12]: 1  ##% derived function
2  def predicted_y(x,w,b):
3      '''
4
5      Parameters
6      -----
7      x : numpy.ndarray
8          DESCRIPTION - array of features (x1...xn)
9      w : numpy.ndarray
10         DESCRIPTION - array of weights (w1...wn)
11      b : float
12         DESCRIPTION - bias
13      Returns
14      -----
15      y : float
16         DESCRIPTION - predicted y based on x
17
18      '''
19      f = np.dot(w,x) + b
20      return f
```

Cost function

In order to judge our prediction error, a cost function is used,

$$J(\vec{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{\vec{w},b}(\vec{x}^i) - y^i)^2$$

This cost function calculates the prediction error - error between the predicted value and the training value.

```
In [13]: 1  # cost function
2  def compute_cost(x_train,y_train,w,b):
3      '''
4      Parameters
5      -----
6      x_train : numpy.ndarray
7          DESCRIPTION - training data features
8      y_train : numpy.ndarray
9          DESCRIPTION : training data targets
10     w : numpy.ndarray
11         DESCRIPTION - array of weights (w1...wn)
12     b : float
13         DESCRIPTION - bias
14     Returns
15     -----
16     y : float
17         DESCRIPTION - predicted y based on x
18
19     '''
20     SUM = 0
21     m = x_train.shape[0]
22     for i in range(m):
23         x = x_train[i]
24         y = y_train[i]
25
26         f = predicted_y(x,w,b)
27
28         SUM += (f-y)**2
29     cost = SUM/(2*m)
30     return cost
```

Compute gradient

Gradient is a partial derivative of the cost function $J(w, b)$ w.r.t. parameters w and b

$$\frac{\partial(J(\vec{w}, b))}{\partial(\vec{w})} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\vec{w},b}(\vec{x}^i) - y^i) \vec{x}^i$$

$$\frac{\partial(J(\vec{w}, b))}{\partial(b)} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\vec{w},b}(\vec{x}^i) - y^i)$$

In [14]: ▶

```
1 # compute gradient
2 def compute_gradient(x_train,y_train,w,b):
3     '''
4     Parameters
5     -----
6     x_train : numpy.ndarray
7         DESCRIPTION - training data features
8     y_train : numpy.ndarray
9         DESCRIPTION : training data targets
10    w : numpy.ndarray
11        DESCRIPTION - array of weights (w1...wn)
12    b : float
13        DESCRIPTION - bias
14
15    Returns
16    -----
17    dj_dw : numpy.ndarray
18        DESCRIPTION - array of gradient values for all parameters
19    dj_db : numpy.float64
20        DESCRIPTION - bias gradient
21    '''
22    SUMw = 0
23    SUMb = 0
24
25    m = x_train.shape[0]
26    SUMw = np.zeros(8)
27    SUMb = 0
28    for i in range(m):
29        x = x_train[i]
30        y = y_train[i]
31
32        f = predicted_y(x,w,b)
33
34        SUMw += np.array([np.dot((f-y),i) for i in x])
35        SUMb += f-y
36
37    dj_dw = SUMw/m
38    dj_db = SUMb/m
39
40    return dj_dw, dj_db
```

Gradient descent

For every iteration the gradient is calculated, the values of w and b needs to be simultaneously updated.

$$\vec{w} = \vec{w} - \alpha \frac{\partial(J(\vec{w}, b)}{\partial(\vec{w})}$$
$$b = b - \alpha \frac{\partial(J(\vec{w}, b)}{\partial(b)}$$

here, α is the learning rate

This step needs to be repetaed until the values of \vec{w} and b converge.

In [15]:

```
1  #gradient_descent
2  def gradient_descent(x_train,y_train,num_iters,alpha,w,b):
3      '''
4      Parameters
5      -----
6      x_train : numpy.ndarray
7          DESCRIPTION - training data features
8      y_train : numpy.ndarray
9          DESCRIPTION : training data targets
10     num_iters : int
11         DESCRIPTION : number of iterations
12     alpha : float
13         DESCRIPTION - learning rate
14     w : numpy.ndarray
15         DESCRIPTION - array of weights (w1...wn)
16     b : float
17         DESCRIPTION - bias
18
19     '''
20
21     J_history = []
22     w_history = []
23     b_history = []
24
25     m = x_train.shape[0]
26
27
28     for i in range(num_iters):
29         dj_dw, dj_db = compute_gradient(x_train,y_train,w,b)
30         w = w - alpha*dj_dw
31         b = b - alpha*dj_db
32
33         cost = compute_cost(x_train, y_train,w,b)
34         J_history.append(cost)
35         w_history.append(w)
36         b_history.append(b)
37
38         if i% math.ceil(num_iters/10) == 0:
39             print(f"Iteration {i:4}: Cost {J_history[-1]:0.2e} ")
40             # f"dj_dw: {dj_dw}, dj_db: {dj_db: 0.3e} ",
41             # f"w: {w}, b:{b: 0.5e}")
42             # break the loop if the cost has converged i.e. the diference in cost is less than 1e-7
43             if i>1 and abs((J_history[-1]) - (J_history[-2])) <= 1e-7:
44                 print('\ncost has converged')
45                 print(f'cost at convergence : {cost}')
46                 break
47     print('w = {}'.format(w,b))
48     return J_history, w_history, b_history #return w and J,w history for graphing
49
```

In [16]:

```
1  # performing multiple linear regression
2  import time
3  tic = time.time()
4  w = np.random.normal(size=8)
5  b = 0
6  num_iters = 10000
7  alpha = 1e-2
8
9  # perform gradient descent
10 J_history, w_history, b_history = gradient_descent(x_train, y_train, num_iters, alpha, w, b)
11 toc = time.time()
12 time = (toc-tic)/60
13 print(f'computation time : {time} minutes')
```


```
Iteration    0: Cost 2.81e+00
Iteration 1000: Cost 5.35e-01
Iteration 2000: Cost 5.34e-01
Iteration 3000: Cost 5.33e-01
```

```
cost has converged
cost at convergence : 0.5334277169272521
w = [ 0.54610662 -0.1116658  -0.34595875  0.84438376 -1.16115836 -0.78686157
      2.00793747 -0.86393735]
b = -0.36799616624792464
computation time : 2.6091659982999165 minutes
```

In [20]:


```
1 #plotting the parameters and cost to check the convergence
2 %matplotlib notebook
3 fig,ax = plt.subplots(10,1, figsize=(10,10))
4 for i in range(8):
5     ax[i].plot(np.array(w_history)[:i,i], label=f'w{i}')
6     ax[i].legend()
7     ax[i].set_ylim((min(np.array(w_history)[:i,i])-1),max(np.array(w_history)[:i,i])+1)
8 ax[8].plot(b_history, label='b')
9 ax[9].plot(J_history, label='cost')
10 plt.legend()
11 plt.tight_layout()
12 plt.show()
```

<IPython.core.display.Javascript object>



```
In [18]: 1 #the parameters and the cost is visibly converged
2
3 #predicting the price using the testing data
4 m = x_test.shape[0]
5 y_pred = []
6
7 for i in range(m):
8     x = x_test[i]
9     y = y_test[i]
10    y_pred.append(predicted_y(x,w,b))
11
12
```

```
In [19]: 1 # plotting the testing and predicted data to check the model
2 %matplotlib inline
3 plt.scatter(y_test, y_pred)
4
5 # Add a diagonal line to the plot
6 plt.plot([min(y_test), max(y_pred)], [min(y_test), max(y_pred)], 'k--', lw=4)
7
8 # Add Labels and title
9 plt.xlabel('Actual Y')
10 plt.ylabel('Predicted Y')
11 plt.title('Comparison of Actual Y vs. Predicted Y')
12
13 # Show the plot
14 plt.show()
```



Summary

The this dataset, using multiple linear regression has definitely improved the prediction instead of just single variable linear regression. The reason behind this improved accuracy is that the use of more than one feature for prediction.

Univariate single linear regression: https://github.com/UmangTyagi/machine_learning_specialization/blob/master/linear_regression/univariate_linear_regression.ipynb
(https://github.com/UmangTyagi/machine_learning_specialization/blob/master/linear_regression/univariate_linear_regression.ipynb)